

# Quantum Field Processor (QFP)

Regulated Execution Architecture for QFM/QVM

Karel Cápek

## Abstract

The Quantum Field Processor (QFP) is the regulated execution core of the QFM/QVM ecosystem. It enforces legitimacy, safety classes, bounded execution, auditability, and compliance across operator workloads, from educational arithmetic kernels to high-risk autonomous orchestration. This paper specifies QFP as a layered architecture: identity and legitimacy gates, permission tokens (NFT-1001), a safety kernel for operator classification and hazard prediction, a sandboxed execution engine, and a zero-knowledge audit and sealing layer suitable for distributed deployments.

## Contents

<b>1</b>	<b>Introduction</b>	<b>5</b>
1.1	From Operator Mathematics to Executable Systems . . . . .	5
1.2	Why Regulation Is a Structural Requirement . . . . .	5
1.3	QFP as an Execution Gate, Not a Compute Core . . . . .	6
1.4	Threat Model and Motivation . . . . .	6
1.5	Core Principles of the QFP Architecture . . . . .	6
1.6	Relation to Quantum-Inspired Computation . . . . .	7
1.7	Scope of This Paper . . . . .	7
1.8	Position Within the QVM Stack . . . . .	7
1.9	Transition to Architectural Specification . . . . .	8
<b>2</b>	<b>Layered Architecture of the Quantum Field Processor</b>	<b>8</b>
2.1	Architectural Overview . . . . .	8
2.2	Layer 1: Identity and Legitimacy . . . . .	8
2.3	Layer 2: Permission and Capability Control . . . . .	9
2.4	Layer 3: Safety Kernel . . . . .	9
2.5	Layer 4: Execution Engine . . . . .	9
2.6	Layer 5: Audit and Result Sealing . . . . .	9
2.7	Separation of Authority . . . . .	10
2.8	Control Flow and Determinism . . . . .	10
2.9	Failure Modes and Safe Termination . . . . .	10
2.10	Extensibility of the Architecture . . . . .	10
2.11	Preparation for Safety Taxonomy . . . . .	11

<b>3 Safety Taxonomy and Operator Classification</b>	<b>11</b>
3.1 Purpose of Safety Classification . . . . .	11
3.2 Operator Programs as Classification Objects . . . . .	11
3.3 Safety Classes Overview . . . . .	11
3.4 Class S1: Benign Operator Programs . . . . .	12
3.5 Class S2: Structured but Limited Programs . . . . .	12
3.6 Class S3: Global Arithmetic Dynamics . . . . .	12
3.7 Class S4: High-Impact or Adaptive Programs . . . . .	13
3.8 Class S5: Prohibited Programs . . . . .	13
3.9 Classification Criteria . . . . .	13
3.10 Conservative Classification Principle . . . . .	13
3.11 Auditability of Classification . . . . .	14
3.12 Transition to Permission Framework . . . . .	14
<b>4 Permission and Capability Framework (NFT-1001)</b>	<b>14</b>
4.1 Rationale for Capability-Based Authorization . . . . .	14
4.2 Definition of an NFT-1001 Certificate . . . . .	14
4.3 Capability Scope and Granularity . . . . .	15
4.4 Binding to Identity . . . . .	15
4.5 Delegation and Sub-Capabilities . . . . .	15
4.6 Temporal Validity and Revocation . . . . .	15
4.7 Emergency Suspension and Kill Switches . . . . .	16
4.8 Verification Within the Execution Pipeline . . . . .	16
4.9 Separation from Safety Classification . . . . .	16
4.10 Auditability of Capabilities . . . . .	16
4.11 Preparation for Execution Pipeline . . . . .	17
<b>5 Execution Pipeline and Control Flow</b>	<b>17</b>
5.1 Execution Request Object . . . . .	17
5.2 Pipeline Overview . . . . .	17
5.3 Phase 1: Request Validation . . . . .	18
5.4 Phase 2: Identity and Legitimacy Verification . . . . .	18
5.5 Phase 3: Capability Verification . . . . .	18
5.6 Phase 4: Safety Analysis and Classification . . . . .	18
5.7 Phase 5: Controlled Execution . . . . .	19
5.8 Phase 6: Result Sealing . . . . .	19
5.9 Phase 7: Audit Commitment . . . . .	19
5.10 Deterministic Control Flow . . . . .	19
5.11 Failure Handling and Safe Termination . . . . .	20
5.12 No Bypass Guarantee . . . . .	20
5.13 Preparation for Distributed Execution . . . . .	20
<b>6 Distributed Shards and Orchestration</b>	<b>20</b>
6.1 Motivation for Sharding . . . . .	20
6.2 Shard Definition . . . . .	20
6.3 Global Versus Local Authority . . . . .	21
6.4 The Orchestrator Role . . . . .	21

6.5	Cross-Shard Execution Requests . . . . .	21
6.6	Prevention of Execution Laundering . . . . .	21
6.7	Global Hazard Accounting . . . . .	22
6.8	Consistency of Safety Decisions . . . . .	22
6.9	Failure Handling in Distributed Contexts . . . . .	22
6.10	Audit and Traceability Across Shards . . . . .	22
6.11	Security Model Summary . . . . .	23
6.12	Preparation for Safety Kernel Algorithms . . . . .	23
<b>7</b>	<b>Safety Kernel Algorithms and Risk Estimation</b>	<b>23</b>
7.1	Design Principles of the Safety Kernel . . . . .	23
7.2	Input Representation . . . . .	23
7.3	Operator Composition Graph Analysis . . . . .	24
7.4	Spectral Radius Estimation . . . . .	24
7.5	Diffusion Cone Analysis . . . . .	24
7.6	Recursion and Feedback Detection . . . . .	24
7.7	External Interface Exposure . . . . .	25
7.8	Composite Risk Score . . . . .	25
7.9	Decision Logic . . . . .	25
7.10	Versioning and Policy Control . . . . .	25
7.11	Auditability and Reproducibility . . . . .	26
7.12	Limitations of Static Analysis . . . . .	26
7.13	Transition to Result Sealing . . . . .	26
<b>8</b>	<b>Result Sealing and Zero-Knowledge Compliance</b>	<b>26</b>
8.1	Purpose of Result Sealing . . . . .	26
8.2	Sealed Result Object . . . . .	27
8.3	Binding Results to Context . . . . .	27
8.4	Confidentiality and Minimal Disclosure . . . . .	27
8.5	Zero-Knowledge Compliance Proofs . . . . .	27
8.6	Compliance Statements . . . . .	28
8.7	Proof Generation and Verification . . . . .	28
8.8	Audit Logs and Append-Only Records . . . . .	28
8.9	Failure and Termination Sealing . . . . .	28
8.10	Interoperability with External Systems . . . . .	29
8.11	Security and Threat Considerations . . . . .	29
8.12	Preparation for Governance Integration . . . . .	29
<b>9</b>	<b>Governance and Regulatory Integration</b>	<b>29</b>
9.1	Separation of Governance and Execution . . . . .	29
9.2	Policy as a Versioned Artifact . . . . .	30
9.3	Policy Activation and Transition . . . . .	30
9.4	Governance Authority Model . . . . .	30
9.5	Scope of Governance Powers . . . . .	30
9.6	Regulatory Interfaces . . . . .	31
9.7	Compliance Without Disclosure . . . . .	31
9.8	Dispute Resolution and Forensic Analysis . . . . .	31

9.9	Emergency Governance Actions . . . . .	31
9.10	Avoidance of Centralized Control . . . . .	32
9.11	Long-Term Governance Evolution . . . . .	32
9.12	Transition to Applications and Future Directions . . . . .	32
<b>10</b>	<b>Applications, Implications, and Future Directions</b>	<b>32</b>
10.1	Application Domains . . . . .	32
10.2	Implications for Computation Models . . . . .	33
10.3	Determinism and Trust . . . . .	33
10.4	Ethical and Societal Considerations . . . . .	33
10.5	Limitations and Open Challenges . . . . .	34
10.6	Future Architectural Extensions . . . . .	34
10.7	Relation to the Broader QVM Stack . . . . .	34
10.8	Concluding Perspective . . . . .	34

# 1 Introduction

The increasing power of computational systems has shifted the primary challenge of computing from raw performance to control, legitimacy, and safety of execution. As computation becomes more autonomous, distributed, and capable of producing globally impactful outcomes, the absence of rigorous execution governance becomes a structural risk rather than a secondary concern.

The Quantum Field Processor (QFP) is introduced as a regulated execution architecture designed to address this problem at a foundational level. Rather than treating safety, permission, and auditability as external policies layered on top of computation, QFP integrates them directly into the execution substrate.

## 1.1 From Operator Mathematics to Executable Systems

The preceding whitepaper (WP-01) established Quansistor Field Mathematics (QFM) as an operator-theoretic framework in which computation is expressed as global evolution over arithmetic Hilbert space. In that setting, operators define computation, spectra encode structure, and execution is interpreted as deterministic operator dynamics.

However, mathematical executability alone is insufficient once such operators are deployed in real systems. When operator programs are capable of:

- large-scale autonomous execution,
- recursive self-composition,
- distributed propagation,
- and real-world actuation,

then unrestricted execution becomes unsafe by default.

QFP exists precisely at this boundary: it is the mechanism by which QFM operators are allowed to execute, restricted, audited, or terminated.

## 1.2 Why Regulation Is a Structural Requirement

In classical computing, regulation is often external:

- operating systems enforce permissions,
- sandboxes isolate processes,
- governance is handled socially or legally.

This separation breaks down for operator-based and quantum-inspired systems. In such systems:

- execution is global rather than local,
- state is distributed rather than compartmentalized,
- intermediate results may already encode irreversible structure.

As a consequence, post-hoc regulation is insufficient. Once an unsafe operator evolution has occurred, its effects may be impossible to undo.

QFP therefore treats regulation as a *pre-execution constraint* rather than a corrective measure.

### 1.3 QFP as an Execution Gate, Not a Compute Core

It is important to emphasize what QFP is not.

QFP:

- is not a mathematical framework (that role belongs to QFM),
- is not a physical processor,
- does not define new operator semantics.

Instead, QFP is an execution gate. It decides:

- whether a requested computation may run,
- under which constraints it may run,
- how it is monitored,
- how its results are sealed and recorded.

All actual computation is still defined at the operator level. QFP governs access to execution.

### 1.4 Threat Model and Motivation

The QFP architecture is motivated by a concrete threat model.

Unregulated operator execution may lead to:

- uncontrolled recursive amplification,
- emergent behavior exceeding design intent,
- covert information exfiltration,
- unsafe coupling between arithmetic operators and external systems,
- inability to audit or reproduce outcomes.

These risks are not hypothetical. They arise naturally in systems that combine:

- global state,
- non-local operator action,
- and spectral feedback mechanisms.

QFP addresses these risks by enforcing a mandatory execution pipeline that every computation must pass through.

### 1.5 Core Principles of the QFP Architecture

The design of QFP is guided by several non-negotiable principles:

**Legitimacy First.** No execution occurs without explicit legitimacy. Every computation request must be attributable, permissioned, and classified before execution.

**Pre-Execution Analysis.** Risk assessment, safety classification, and boundary checking occur before any operator is applied.

**Deterministic Enforcement.** All regulatory decisions are deterministic, reproducible, and auditable. There are no opaque heuristic gates.

**Separation of Power.** No single component may both authorize and execute unrestricted operator flows.

**Auditability by Construction.** Execution results are sealed together with metadata sufficient for later verification, without revealing sensitive internal state.

## 1.6 Relation to Quantum-Inspired Computation

Although QFP is designed to support quantum-inspired operator dynamics, it does not rely on probabilistic measurement, randomness, or physical quantum effects.

Instead:

- execution remains deterministic,
- inspection remains possible,
- rollback and termination are defined operations.

This makes QFP suitable for regulated environments where reproducibility and accountability are mandatory.

## 1.7 Scope of This Paper

This paper specifies QFP as an abstract execution architecture independent of any particular deployment platform.

It will:

- define the layered structure of QFP,
- formalize safety classes and execution permissions,
- describe the execution pipeline from request to result sealing,
- outline mechanisms for distributed orchestration and governance.

It will not:

- prescribe a specific hardware implementation,
- mandate a particular blockchain or network,
- replace mathematical analysis performed at the QFM level.

## 1.8 Position Within the QVM Stack

Within the broader QVM ecosystem:

- QFM defines what computation *is*,
- QFP defines whether and how it may execute,
- QVM defines where and at what scale it runs.

This separation ensures that mathematical expressiveness, execution safety, and architectural scalability evolve independently but coherently.

## 1.9 Transition to Architectural Specification

With the motivation and role of QFP established, we now turn to its internal structure. The next section presents the layered architecture of the Quantum Field Processor and explains how legitimacy, safety analysis, execution, and audit are composed into a single deterministic pipeline.

# 2 Layered Architecture of the Quantum Field Processor

The Quantum Field Processor (QFP) is defined as a strictly layered execution architecture. Each layer performs a well-defined function and communicates with adjacent layers through explicit, verifiable interfaces. No layer is permitted to bypass another, and no single layer possesses sufficient authority to authorize and execute unrestricted computation on its own.

This section describes the architectural layers of QFP, their roles, and their interaction within the regulated execution pipeline.

## 2.1 Architectural Overview

At a high level, QFP is composed of five primary layers:

1. Identity and Legitimacy Layer
2. Permission and Capability Layer
3. Safety Kernel
4. Execution Engine
5. Audit and Result Sealing Layer

Each execution request traverses these layers sequentially. Failure or rejection at any layer terminates the request deterministically and produces an auditable decision record.

## 2.2 Layer 1: Identity and Legitimacy

The first layer establishes *who* is requesting execution and under which identity the request is made.

This layer performs:

- identity verification,
- attribution of responsibility,
- association with a legitimacy context.

Identity is treated abstractly: it may correspond to a human operator, an organization, an automated agent, or a delegated system. QFP does not assume a specific identity technology, but requires that identity claims be cryptographically verifiable and non-forgeable.

No execution proceeds without a validated identity context.

## 2.3 Layer 2: Permission and Capability Control

Once identity is established, QFP evaluates whether the requester possesses the necessary permissions to invoke the requested computation.

Permissions are expressed as explicit capability certificates rather than implicit roles. These certificates define:

- which classes of operators may be executed,
- allowable resource bounds,
- recursion and composition limits,
- temporal and contextual constraints.

This layer enforces the principle of least authority. Possession of a general execution capability does not imply permission to execute high-impact or self-modifying operator programs.

## 2.4 Layer 3: Safety Kernel

The Safety Kernel is the core analytical component of QFP. It performs pre-execution analysis of the requested computation and determines whether execution is admissible under the current policy.

Its responsibilities include:

- classification of operator programs into safety classes,
- detection of forbidden compositions,
- estimation of dynamic risk indicators,
- enforcement of execution boundaries.

The Safety Kernel operates deterministically and produces an explicit decision: allow, restrict, sandbox, or deny. Its decision logic is auditable and versioned.

## 2.5 Layer 4: Execution Engine

Only after passing all prior layers does the request reach the Execution Engine.

The Execution Engine:

- executes operator programs under enforced constraints,
- monitors resource usage and execution invariants,
- supports interruption and termination,
- isolates execution contexts.

Importantly, the Execution Engine has no authority to relax permissions or override safety decisions. It executes strictly within the envelope defined by earlier layers.

## 2.6 Layer 5: Audit and Result Sealing

The final layer is responsible for producing verifiable execution records.

This layer:

- seals execution results together with metadata,

- records policy decisions and safety classifications,
- produces cryptographic commitments suitable for later verification,
- supports compliance and forensic analysis.

Result sealing ensures that outputs cannot be separated from the context in which they were produced.

## 2.7 Separation of Authority

A fundamental design constraint of QFP is strict separation of authority.

- No layer both authorizes and executes computation.
- No execution occurs without prior safety analysis.
- No result exists without an audit trail.

This separation prevents single-point failures and reduces the risk of escalation through component compromise.

## 2.8 Control Flow and Determinism

The layered architecture enforces deterministic control flow.

Given:

- the same execution request,
- the same policy state,
- the same Safety Kernel version,

the outcome of the pipeline is identical.

This property is essential for reproducibility, governance, and legal accountability.

## 2.9 Failure Modes and Safe Termination

Each layer defines explicit failure modes.

- Identity failure produces immediate rejection.
- Permission failure produces denial with explanation.
- Safety failure produces classified rejection or sandboxing.
- Execution failure produces a sealed termination record.

No failure mode results in undefined or partially executed computation.

## 2.10 Extensibility of the Architecture

The layered structure is intentionally extensible.

New safety checks, permission types, or audit mechanisms may be introduced by extending individual layers without modifying the entire system. This supports long-term evolution without architectural breakage.

## 2.11 Preparation for Safety Taxonomy

With the layered architecture defined, we now turn to the internal logic of the Safety Kernel. The next section introduces a formal taxonomy of safety classes and explains how operator programs are classified and constrained prior to execution.

# 3 Safety Taxonomy and Operator Classification

Safe execution within the Quantum Field Processor requires a precise and deterministic classification of operator programs prior to execution. Because QFM operators act globally, admit composition, and may exhibit emergent behavior, safety cannot be inferred from individual operators alone. It must be assessed at the level of operator programs and their anticipated dynamic effects.

This section introduces the QFP safety taxonomy, defines safety classes, and specifies how operator programs are classified and constrained by the Safety Kernel.

### 3.1 Purpose of Safety Classification

Safety classification serves three essential purposes:

- to prevent execution of inherently unsafe operator programs,
- to bound the behavior of partially safe programs,
- to provide a deterministic and auditable decision basis.

Classification occurs *before execution* and is a mandatory gate in the QFP pipeline. No operator program may execute without an assigned safety class.

### 3.2 Operator Programs as Classification Objects

The object of classification is not an individual operator, but an *operator program*, defined as:

A finite or parameterized composition of QFM operators, together with declared execution parameters, potential terms, recursion bounds, and external interfaces.

Safety is therefore a property of:

- operator composition structure,
- parameter ranges,
- allowed iteration or evolution time,
- interaction with external systems.

### 3.3 Safety Classes Overview

QFP defines five ordered safety classes, denoted S1 through S5. These classes are monotonic: higher classes strictly dominate lower ones in terms of potential impact and required controls.

Class	Informal Description
S1	Benign, bounded, non-amplifying execution
S2	Structured computation with limited global effects
S3	Global arithmetic dynamics under strict bounds
S4	High-impact or adaptive operator programs
S5	Prohibited or uncontrollable execution

### 3.4 Class S1: Benign Operator Programs

Class S1 includes operator programs that are provably bounded and non-amplifying.

Typical properties:

- finite support or explicitly bounded state space,
- no recursion or self-composition,
- no external interfaces,
- guaranteed termination.

S1 programs may be executed with minimal restrictions and are suitable for educational, diagnostic, or low-risk computational tasks.

### 3.5 Class S2: Structured but Limited Programs

Class S2 includes operator programs that act on larger arithmetic domains but remain structurally constrained.

Typical properties:

- bounded diffusion or evolution time,
- restricted operator sets,
- no adaptive modification of parameters,
- predictable spectral envelope.

S2 execution requires resource monitoring but does not require sandboxing or external isolation.

### 3.6 Class S3: Global Arithmetic Dynamics

Class S3 covers operator programs that generate genuine global arithmetic dynamics.

Typical properties:

- unbounded index space,
- significant spectral interference,
- long-range multiplicative propagation,
- sensitivity to parameter tuning.

S3 programs are permitted only under strict execution bounds, mandatory monitoring, and enforced termination conditions.

### **3.7 Class S4: High-Impact or Adaptive Programs**

Class S4 includes operator programs whose behavior may change during execution or whose outputs may influence future executions.

Examples include:

- adaptive Hamiltonians,
- feedback-driven parameter updates,
- coupling to external control systems,
- recursive self-application.

S4 execution is highly restricted and requires explicit authorization, sandboxing, and enhanced audit guarantees.

### **3.8 Class S5: Prohibited Programs**

Class S5 denotes operator programs that are disallowed under all standard conditions.

This includes:

- unbounded self-referential execution,
- attempts to bypass Safety Kernel constraints,
- programs with undefined termination behavior,
- operators designed to evade audit or attribution.

S5 classification results in immediate rejection and permanent logging of the attempt.

### **3.9 Classification Criteria**

The Safety Kernel assigns a safety class by evaluating a fixed set of criteria, including:

- operator composition graph,
- estimated spectral radius,
- diffusion cone size,
- recursion depth and feedback loops,
- external interface exposure.

Each criterion is evaluated deterministically and contributes to the final classification decision.

### **3.10 Conservative Classification Principle**

When uncertainty exists, the Safety Kernel applies a conservative principle:

If a program cannot be proven safe within a given class, it is classified into the next higher risk class.

This ensures that classification errors err on the side of restriction rather than permissiveness.

### **3.11 Auditability of Classification**

Every classification decision produces an auditable record containing:

- the assigned safety class,
- the evaluated criteria,
- the Safety Kernel version identifier,
- the decision timestamp.

These records are immutable and may be reviewed independently of execution results.

### **3.12 Transition to Permission Framework**

Safety classification determines *what class* of execution is requested, but not *whether* the requester is authorized to perform it.

The next section introduces the permission and capability framework that governs access to safety classes and binds execution rights to explicit certificates.

## **4 Permission and Capability Framework (NFT-1001)**

Safety classification determines the inherent risk of an operator program, but it does not determine who is authorized to execute it. The Quantum Field Processor therefore separates safety assessment from execution authority.

This section introduces the permission and capability framework used by QFP. Execution rights are represented as explicit, verifiable capability certificates, collectively referred to as the NFT-1001 framework.

### **4.1 Rationale for Capability-Based Authorization**

Traditional role-based authorization models are insufficient for operator-based computation. Roles are coarse, implicit, and difficult to audit. In contrast, QFP adopts a capability-based model in which authority is expressed as an explicit, transferable, and revocable object.

A capability certificate answers a precise question:

Which classes of operator programs may this identity execute, under which constraints, and for how long?

NFT-1001 is the canonical representation of such a capability.

### **4.2 Definition of an NFT-1001 Certificate**

An NFT-1001 certificate is defined abstractly as a cryptographically verifiable data object containing execution permissions.

Formally, an NFT-1001 certificate encodes:

- the issuing authority,
- the subject identity or delegation scope,
- the permitted safety classes,
- execution bounds and constraints,

- validity and revocation conditions.

The term “NFT” is used descriptively to indicate non-fungibility and unique identity; the framework does not mandate a specific token standard or ledger technology.

### **4.3 Capability Scope and Granularity**

Execution permissions are intentionally granular.

An NFT-1001 certificate may specify:

- a maximum safety class (e.g., S2 or S3),
- allowable operator families,
- bounds on execution time or iteration count,
- restrictions on external interfaces,
- limits on recursion or self-composition.

Possession of a lower-class capability does not imply permission to execute higher-class programs.

### **4.4 Binding to Identity**

Each NFT-1001 certificate is bound to an identity context.

Binding may take one of the following forms:

- direct binding to a single identity,
- delegation to a bounded identity set,
- conditional binding based on execution context.

QFP requires that identity binding be verifiable and non-repudiable. No anonymous execution of privileged operator programs is permitted.

### **4.5 Delegation and Sub-Capabilities**

NFT-1001 certificates may support controlled delegation.

A certificate holder may issue a sub-capability with:

- strictly reduced privileges,
- shorter validity,
- narrower execution scope.

Delegation chains are explicitly represented and auditable. Cyclic or unbounded delegation is prohibited by construction.

### **4.6 Temporal Validity and Revocation**

Capabilities are not perpetual by default.

Each NFT-1001 certificate includes:

- an explicit validity interval,
- revocation conditions,

- optional emergency termination flags.

Revocation is treated as a first-class operation. Once revoked, a capability cannot be used to authorize execution, even if cached by an execution node.

#### 4.7 Emergency Suspension and Kill Switches

For high-impact safety classes (S3 and above), QFP supports emergency suspension mechanisms.

These include:

- immediate revocation of active capabilities,
- forced termination of running executions,
- global invalidation of specific capability classes.

Emergency actions are auditable and require explicit authorization under governance policy.

#### 4.8 Verification Within the Execution Pipeline

During execution request processing, QFP verifies:

- the authenticity of the NFT-1001 certificate,
- its binding to the requester identity,
- compatibility with the assigned safety class,
- compliance with all encoded constraints.

Failure at any verification step results in deterministic denial of execution.

#### 4.9 Separation from Safety Classification

A critical design principle is strict separation between:

- *safety classification* (what is inherently risky),
- *permission authorization* (who is allowed to take that risk).

The Safety Kernel does not grant permission, and the capability framework does not override safety decisions. Both must independently approve execution.

#### 4.10 Auditability of Capabilities

Every use of an NFT-1001 certificate produces an audit record containing:

- certificate identifier,
- identity binding,
- permitted safety class,
- execution context,
- verification outcome.

This enables post-hoc analysis of who was authorized to do what, and under which conditions.

## 4.11 Preparation for Execution Pipeline

With safety classification and permission authorization defined, the conditions for execution are now fully specified.

The next section describes the execution pipeline itself, tracing the lifecycle of an execution request from submission through authorization, analysis, execution, and result sealing.

# 5 Execution Pipeline and Control Flow

The Quantum Field Processor enforces execution through a strictly defined, deterministic pipeline. Every execution request follows the same sequence of phases, with explicit inputs, outputs, and failure modes. No phase may be skipped, reordered, or partially executed.

This section specifies the execution pipeline of QFP and describes the control flow governing request handling, decision points, and termination.

## 5.1 Execution Request Object

Execution begins with the submission of an *execution request*. This request is a structured object containing all information required for authorization, analysis, execution, and audit.

An execution request includes:

- requester identity reference,
- operator program specification,
- declared parameters and bounds,
- requested execution context,
- capability certificate reference (NFT-1001).

Requests lacking required fields are rejected prior to pipeline entry.

## 5.2 Pipeline Overview

The QFP execution pipeline consists of seven sequential phases:

1. Request Validation
2. Identity and Legitimacy Verification
3. Capability Verification
4. Safety Analysis and Classification
5. Controlled Execution
6. Result Sealing
7. Audit Commitment

Each phase produces a well-defined output that becomes the input of the next phase.

### **5.3 Phase 1: Request Validation**

The initial phase verifies syntactic and structural correctness.

This includes:

- schema validation of the request object,
- completeness of declared parameters,
- consistency between requested operators and declared bounds.

Failure at this stage results in immediate rejection without side effects.

### **5.4 Phase 2: Identity and Legitimacy Verification**

In this phase, QFP verifies that the request is associated with a valid and verifiable identity.

Operations include:

- cryptographic verification of identity claims,
- attribution of responsibility,
- binding of identity to the execution context.

If identity verification fails, the request is denied and logged.

### **5.5 Phase 3: Capability Verification**

The capability certificate referenced in the request is verified.

This phase checks:

- authenticity and integrity of the NFT-1001 certificate,
- binding between certificate and requester identity,
- validity interval and revocation status,
- compatibility with requested execution scope.

Requests exceeding capability scope are deterministically denied.

### **5.6 Phase 4: Safety Analysis and Classification**

The Safety Kernel evaluates the operator program.

This phase performs:

- operator composition analysis,
- estimation of spectral and diffusion indicators,
- assignment of a safety class,
- determination of execution constraints.

The outcome is one of the following:

- approval with constraints,
- approval with sandboxing,
- rejection.

## 5.7 Phase 5: Controlled Execution

Approved requests enter the Execution Engine.

Execution proceeds under enforced constraints:

- resource bounds,
- recursion and iteration limits,
- isolation requirements,
- monitoring hooks.

The Execution Engine continuously enforces invariants and may terminate execution if violations are detected.

## 5.8 Phase 6: Result Sealing

Upon normal termination or forced interruption, execution results are sealed.

Sealing includes:

- hashing of outputs,
- binding to execution metadata,
- inclusion of safety and capability identifiers.

Sealed results are immutable and verifiable.

## 5.9 Phase 7: Audit Commitment

The final phase records an audit commitment.

This includes:

- execution outcome,
- timestamps and identifiers,
- decision traces from prior phases.

Audit commitments are stored in an append-only log suitable for independent verification.

## 5.10 Deterministic Control Flow

At every phase, control flow is deterministic.

Given:

- identical requests,
- identical policy and kernel versions,
- identical capability states,

the pipeline produces identical outcomes.

This property is essential for reproducibility and governance.

## 5.11 Failure Handling and Safe Termination

Failures at any phase result in explicit termination.

No partial execution occurs unless all prior phases have completed successfully. Forced termination during execution produces a sealed termination record indicating the cause and execution state.

## 5.12 No Bypass Guarantee

The pipeline enforces a no-bypass guarantee.

- No execution may occur without passing all prior phases.
- No internal component may inject execution directly.
- No cached decision may override real-time verification.

This guarantee is enforced by architectural separation and explicit interfaces.

## 5.13 Preparation for Distributed Execution

The execution pipeline described here is defined abstractly and applies equally to single-node and distributed deployments.

The next section extends this pipeline to distributed and sharded environments, introducing orchestration, coordination, and cross-shard safety enforcement.

# 6 Distributed Shards and Orchestration

Large-scale deployment of the Quantum Field Processor requires distribution across multiple execution nodes. Distribution introduces scalability, but also amplifies risk: unsafe behavior may emerge not from a single execution context, but from coordinated or sequential actions across shards.

This section defines the distributed execution model of QFP, the role of the orchestrator, and the mechanisms that preserve global safety invariants across sharded environments.

## 6.1 Motivation for Sharding

Sharding is introduced for practical reasons:

- scalability of execution throughput,
- parallel evaluation of operator programs,
- fault isolation and resilience,
- geographic or jurisdictional separation.

However, sharding must not weaken safety guarantees. QFP treats distribution as a performance optimization, not as a relaxation of control semantics.

## 6.2 Shard Definition

A shard is defined as an execution domain capable of:

- running the full QFP execution pipeline,

- enforcing local safety and permission decisions,
- producing sealed execution results.

Each shard is autonomous in execution but subordinate to global policy and orchestration constraints.

### 6.3 Global Versus Local Authority

QFP enforces a strict separation between local and global authority.

- Local shards may execute only within explicitly granted scopes.
- Global authority defines safety policy, capability rules, and revocation signals.
- No shard may redefine safety classes or permission semantics.

This prevents policy drift and inconsistent enforcement.

### 6.4 The Orchestrator Role

The orchestrator is a coordination component responsible for maintaining global execution coherence.

Its responsibilities include:

- routing execution requests to shards,
- enforcing global execution limits,
- tracking distributed execution state,
- aggregating audit commitments.

The orchestrator does not execute operator programs itself. It governs where and how execution occurs.

### 6.5 Cross-Shard Execution Requests

Some operator programs require execution across multiple shards.

For such cases, QFP requires:

- explicit declaration of cross-shard execution,
- pre-allocation of shard scopes,
- unified safety classification applied globally.

Partial approval is not permitted. Either all involved shards approve execution under identical constraints, or execution is denied.

### 6.6 Prevention of Execution Laundering

Execution laundering refers to the attempt to decompose a prohibited or high-risk program into multiple lower-risk executions across shards.

QFP prevents this through:

- global tracking of execution lineage,
- aggregation of safety impact across shards,

- detection of correlated execution patterns.

The orchestrator maintains a hazard map that records cumulative execution effects and enforces escalation when thresholds are crossed.

## 6.7 Global Hazard Accounting

Each shard reports execution metrics to the orchestrator, including:

- spectral impact estimates,
- diffusion cone growth,
- recursion and feedback indicators.

The orchestrator aggregates these metrics to assess global risk. Execution that appears benign locally may be halted if its global footprint becomes unsafe.

## 6.8 Consistency of Safety Decisions

To ensure consistent classification:

- Safety Kernel versions are globally pinned,
- classification rules are immutable during execution,
- updates require coordinated policy transitions.

This guarantees that identical operator programs receive identical safety decisions across shards.

## 6.9 Failure Handling in Distributed Contexts

Distributed execution introduces additional failure modes:

- shard unavailability,
- partial execution completion,
- network partitioning.

QFP mandates fail-safe behavior:

- execution halts on loss of orchestration coherence,
- partial results are sealed and marked incomplete,
- re-execution requires fresh authorization.

No silent continuation is permitted under degraded conditions.

## 6.10 Audit and Traceability Across Shards

Every shard produces local audit records. The orchestrator aggregates these into a global execution trace.

This trace:

- preserves execution ordering,
- binds shard-local results into a single record,

- enables post-hoc reconstruction of distributed behavior.

Audit aggregation does not require access to raw execution data, preserving confidentiality.

## 6.11 Security Model Summary

The distributed QFP model ensures that:

- sharding improves scalability without reducing safety,
- no execution escapes global oversight,
- distributed behavior remains auditable and accountable.

Distribution is thus a controlled extension of QFP semantics, not an exception.

## 6.12 Preparation for Safety Kernel Algorithms

With distributed execution defined, we now turn to the internal mechanics of the Safety Kernel itself.

The next section specifies the analytical algorithms used to classify, bound, and constrain operator programs prior to execution.

# 7 Safety Kernel Algorithms and Risk Estimation

The Safety Kernel is the analytical core of the Quantum Field Processor. Its task is to evaluate operator programs *prior to execution* and to produce a deterministic safety decision based on formal, reproducible criteria.

This section specifies the internal algorithms of the Safety Kernel, the risk indicators it computes, and the decision logic by which operator programs are classified and constrained.

## 7.1 Design Principles of the Safety Kernel

The Safety Kernel is governed by the following principles:

- determinism: identical inputs yield identical outputs,
- conservatism: uncertainty increases restriction, not permissiveness,
- transparency: all decisions are explainable and auditable,
- non-bypassability: no execution proceeds without kernel approval.

The kernel does not learn, adapt, or optimize itself. All algorithms are static, versioned, and explicitly defined.

## 7.2 Input Representation

The Safety Kernel operates on a normalized internal representation of the execution request, consisting of:

- operator composition graph,
- declared parameters and bounds,
- requested evolution time or iteration count,

- declared external interfaces,
- execution context metadata.

This representation is canonicalized to prevent semantic ambiguity.

### 7.3 Operator Composition Graph Analysis

Operator programs are represented as directed graphs whose nodes are operators and whose edges represent composition, iteration, or feedback.

The kernel analyzes:

- graph depth and breadth,
- presence of cycles or self-references,
- unbounded composition paths.

Graphs exceeding predefined structural limits are immediately escalated to higher safety classes or rejected.

### 7.4 Spectral Radius Estimation

A central risk indicator is the estimated spectral radius of the induced operator.

The kernel computes an upper bound on the spectral radius using:

- operator norm estimates of individual components,
- composition amplification bounds,
- worst-case parameter ranges.

If the estimated spectral radius exceeds class-specific thresholds, execution is restricted or denied.

### 7.5 Diffusion Cone Analysis

Diffusion cone analysis estimates the rate at which arithmetic amplitudes spread across index space.

The kernel evaluates:

- expected growth of support over time,
- multiplicative branching factors,
- interaction across prime channels.

Rapidly expanding diffusion cones indicate potential for uncontrolled global effects and result in safety escalation.

### 7.6 Recursion and Feedback Detection

Recursive operator application is a major source of emergent behavior.

The kernel detects:

- explicit recursion constructs,
- implicit feedback via state-dependent parameters,

- chained executions forming effective loops.

Recursion depth and feedback strength are bounded explicitly. Unbounded or opaque recursion results in classification as S4 or S5.

## 7.7 External Interface Exposure

Operator programs may interact with external systems through declared interfaces.

The kernel evaluates:

- directionality of data flow,
- potential for control coupling,
- amplification of external signals.

Programs with write-capable or bidirectional interfaces are subject to stricter safety thresholds.

## 7.8 Composite Risk Score

Individual risk indicators are combined into a composite risk score.

Combination rules are deterministic and monotonic:

- no indicator may reduce overall risk,
- exceeding any critical threshold escalates classification,
- uncertainty contributes positively to risk.

The composite score is mapped to a safety class according to fixed policy tables.

## 7.9 Decision Logic

The Safety Kernel produces one of four decisions:

- approve,
- approve with constraints,
- approve with sandboxing,
- deny.

Each decision is accompanied by:

- assigned safety class,
- enforced execution bounds,
- justification summary.

## 7.10 Versioning and Policy Control

All Safety Kernel algorithms are versioned.

Execution decisions record:

- kernel version identifier,
- policy configuration hash,

- decision timestamp.

Policy updates require explicit governance action and do not retroactively alter past decisions.

### 7.11 Auditability and Reproducibility

Given the same execution request and policy version, the Safety Kernel must produce identical outputs.

This property allows:

- independent verification of decisions,
- forensic analysis after incidents,
- reproducible compliance reporting.

### 7.12 Limitations of Static Analysis

The Safety Kernel does not claim perfect prediction of all dynamic behavior.

Instead:

- it enforces conservative bounds,
- it limits worst-case impact,
- it ensures that unsafe behavior cannot escalate silently.

Residual risk is managed through execution constraints and audit, not optimistic prediction.

### 7.13 Transition to Result Sealing

With execution bounded and monitored, the final concern is integrity and verifiability of results.

The next section describes how execution outputs are sealed, committed, and made verifiable through cryptographic and zero-knowledge mechanisms.

## 8 Result Sealing and Zero-Knowledge Compliance

Regulated execution is incomplete unless its outcomes are verifiable, auditable, and attributable without compromising confidentiality. The Quantum Field Processor therefore treats result sealing and compliance as first-class components of the execution pipeline.

This section defines how execution results are sealed, how compliance is demonstrated without disclosure, and how verifiable records are produced for governance and external review.

### 8.1 Purpose of Result Sealing

Result sealing serves multiple objectives simultaneously:

- integrity of execution outputs,
- binding of results to execution context,

- prevention of tampering or selective disclosure,
- support for later verification and compliance.

A sealed result is not merely an output value, but a cryptographically committed execution artifact.

## 8.2 Sealed Result Object

Upon completion or termination of execution, QFP produces a *sealed result object*.

This object contains:

- a cryptographic commitment to execution outputs,
- a hash of the operator program specification,
- identifiers of applied safety class and capability certificate,
- execution bounds and termination status,
- kernel and policy version identifiers.

The sealed result object is immutable once created.

## 8.3 Binding Results to Context

A central requirement is that results cannot be separated from the context in which they were produced.

Binding ensures that:

- outputs cannot be reused under different permissions,
- results cannot be misattributed to different identities,
- partial results cannot be selectively extracted.

This is achieved by including all relevant contextual identifiers in the commitment hash.

## 8.4 Confidentiality and Minimal Disclosure

QFP explicitly separates:

- verification of correctness and compliance,
- disclosure of execution content.

Sealed results do not require publication of:

- internal state trajectories,
- intermediate operator amplitudes,
- sensitive parameters or inputs.

Only cryptographic commitments and metadata necessary for verification are exposed.

## 8.5 Zero-Knowledge Compliance Proofs

To enable verification without disclosure, QFP supports zero-knowledge compliance proofs.

Such proofs demonstrate statements of the form:

This execution was authorized, classified, and bounded according to policy, and its result was produced by an execution conforming to those bounds.

The proof reveals no additional information beyond the truth of the statement itself.

## 8.6 Compliance Statements

Typical compliance statements include:

- the execution did not exceed declared resource limits,
- the operator program belonged to an approved safety class,
- no prohibited external interfaces were accessed,
- execution terminated within allowed bounds.

Each statement corresponds to a formally defined predicate over the execution record.

## 8.7 Proof Generation and Verification

Proof generation occurs at sealing time and is bound to the sealed result object.

Verification may be performed by:

- internal governance components,
- external auditors,
- regulatory or compliance authorities.

Verification does not require trust in the execution node beyond standard cryptographic assumptions.

## 8.8 Audit Logs and Append-Only Records

Sealed results are referenced in append-only audit logs.

Audit logs record:

- execution identifiers,
- sealed result commitments,
- timestamps and ordering information,
- policy and kernel version references.

Logs may be replicated, archived, or notarized without exposing execution content.

## 8.9 Failure and Termination Sealing

Abnormal termination is treated as a valid outcome and is sealed accordingly.

Termination seals include:

- reason for termination,
- execution state at termination,
- applied enforcement action.

This prevents silent failure and ensures accountability even for aborted executions.

## 8.10 Interoperability with External Systems

The sealing framework is designed to be interoperable.

Sealed results and proofs may be:

- exported to external compliance systems,
- anchored to independent archival services,
- referenced by governance or dispute-resolution processes.

No single external platform is assumed or required.

## 8.11 Security and Threat Considerations

The sealing and proof mechanisms assume:

- standard cryptographic hash security,
- soundness of zero-knowledge constructions,
- integrity of policy versioning.

Attacks attempting to forge compliance without valid execution are detectable through verification failure.

## 8.12 Preparation for Governance Integration

With result integrity and compliance verification established, the final architectural concern is governance.

The next section describes how policy evolution, oversight, and regulatory interfaces are integrated into QFP without undermining deterministic execution guarantees.

# 9 Governance and Regulatory Integration

While execution within the Quantum Field Processor is strictly deterministic, the rules governing which executions are permitted are not static. Safety thresholds, permission semantics, and compliance criteria must evolve in response to new knowledge, emerging risks, and societal constraints.

This section describes how governance and regulatory oversight are integrated into QFP without compromising determinism, auditability, or technical integrity.

## 9.1 Separation of Governance and Execution

A foundational principle of QFP is the separation between:

- *governance*, which defines policy and rules,
- *execution*, which applies those rules deterministically.

Governance mechanisms may change policies, but they do not influence the outcome of individual executions retroactively. Execution nodes apply only the policy version active at request time.

## 9.2 Policy as a Versioned Artifact

All governance decisions are expressed as explicit, versioned policy artifacts.

A policy artifact specifies:

- safety class thresholds,
- permission constraints and capability semantics,
- Safety Kernel configuration parameters,
- compliance predicates and audit requirements.

Each artifact is immutable once published and identified by a unique version hash.

## 9.3 Policy Activation and Transition

Policy updates are not applied instantaneously or implicitly.

QFP enforces:

- explicit activation timestamps,
- forward-only applicability,
- coexistence of multiple policy versions during transition periods.

Execution requests are evaluated under the policy version active at submission time. This guarantees reproducibility and legal clarity.

## 9.4 Governance Authority Model

QFP does not mandate a single governance structure.

Possible governance authorities include:

- institutional review boards,
- multi-stakeholder councils,
- regulated organizational entities,
- formally specified automated governance processes.

What matters is not who governs, but that governance actions are cryptographically attributable and auditable.

## 9.5 Scope of Governance Powers

Governance may:

- define or modify safety class thresholds,
- issue or revoke capability authorities,
- authorize emergency suspension mechanisms,
- mandate additional compliance predicates.

Governance may not:

- override Safety Kernel decisions for individual executions,
- bypass the execution pipeline,

- alter sealed historical records.

This ensures that governance cannot arbitrarily intervene in execution.

## 9.6 Regulatory Interfaces

QFP is designed to interface with external regulatory frameworks without embedding jurisdiction-specific logic.

Regulatory integration is achieved through:

- policy artifacts encoding regulatory constraints,
- compliance predicates reflecting legal requirements,
- verifiable audit records suitable for inspection.

Jurisdictional differences are expressed as policy differences, not code branches.

## 9.7 Compliance Without Disclosure

A central requirement of regulatory integration is minimal disclosure.

QFP enables regulators to verify:

- that an execution was authorized,
- that safety bounds were respected,
- that results were sealed correctly,

without granting access to sensitive execution content or intellectual property.

This is achieved through sealed results and zero-knowledge compliance proofs.

## 9.8 Dispute Resolution and Forensic Analysis

In the event of disputes or incidents, QFP provides a forensic trail.

Forensic analysis may reconstruct:

- the exact policy version applied,
- the Safety Kernel decision logic,
- the execution bounds and termination conditions.

This reconstruction does not require re-execution of the computation and does not rely on trust in the executing party.

## 9.9 Emergency Governance Actions

QFP supports emergency governance actions for exceptional circumstances.

Such actions include:

- immediate revocation of capability classes,
- suspension of execution categories,
- enforced termination of active executions.

Emergency actions are logged, attributable, and subject to post-hoc review.

## 9.10 Avoidance of Centralized Control

Although governance authority exists, QFP avoids centralized technical control.

Execution nodes do not trust governance blindly; they verify policy artifacts cryptographically and enforce them locally. This prevents covert policy manipulation and single-point compromise.

## 9.11 Long-Term Governance Evolution

The governance model of QFP is intentionally extensible.

As new risks, legal standards, or execution paradigms emerge, governance can adapt by issuing new policy artifacts without modifying the execution engine or Safety Kernel algorithms.

## 9.12 Transition to Applications and Future Directions

With governance integration established, the QFP architecture is complete.

The final section discusses applications, broader implications, and possible future directions enabled by regulated operator execution.

# 10 Applications, Implications, and Future Directions

The Quantum Field Processor has been specified as a regulated execution architecture capable of enforcing safety, legitimacy, and auditability for operator-based computation. This section discusses practical application domains, broader systemic implications, and realistic directions for future development.

The emphasis is on applicability and responsibility rather than speculation.

## 10.1 Application Domains

The QFP architecture is applicable wherever computation exhibits global effects, non-local propagation, or high-impact outcomes.

**Scientific and Mathematical Computation.** QFP enables safe execution of operator-heavy workloads, including:

- large-scale spectral analysis,
- arithmetic and number-theoretic experimentation,
- simulation of operator dynamics with global coupling.

Safety classification ensures that exploratory computation remains bounded and reproducible.

**Regulated Research Environments.** Institutions conducting sensitive research may deploy QFP to:

- enforce execution policies,
- restrict high-impact computation to authorized identities,
- maintain verifiable audit trails.

**Autonomous and Semi-Autonomous Systems.** Systems that adapt or act autonomously can use QFP to ensure that:

- execution privileges are explicitly granted,
- adaptive behavior remains within declared bounds,
- accountability is preserved despite autonomy.

**Distributed Computational Infrastructure.** QFP supports safe deployment across distributed nodes, enabling:

- parallel operator execution,
- cross-shard coordination,
- jurisdiction-aware policy enforcement.

## 10.2 Implications for Computation Models

QFP implies a shift in how computation is governed.

Rather than treating safety and compliance as external concerns, QFP embeds them into the execution substrate itself. This changes:

- how computation is authorized,
- how responsibility is assigned,
- how outcomes are validated.

Execution becomes a regulated act, not merely a technical operation.

## 10.3 Determinism and Trust

By enforcing deterministic execution and verifiable audit records, QFP reduces reliance on institutional trust.

Trust is replaced by:

- explicit policy artifacts,
- cryptographic verification,
- reproducible decision logic.

This is particularly relevant in environments where multiple stakeholders with differing incentives interact.

## 10.4 Ethical and Societal Considerations

QFP does not attempt to encode ethics directly. Instead, it provides a mechanism by which ethical constraints may be enforced operationally.

This distinction is deliberate:

- ethical values are expressed through governance policy,
- enforcement is handled deterministically by the system,
- disagreements remain visible and auditable.

By making policy explicit, QFP avoids hidden normative assumptions.

## 10.5 Limitations and Open Challenges

The architecture presented here has explicit limitations.

- Safety analysis is conservative and may restrict benign programs.
- Static analysis cannot predict all emergent behavior.
- Governance processes introduce procedural overhead.

These limitations are accepted trade-offs in exchange for bounded risk and accountability.

## 10.6 Future Architectural Extensions

Several extensions are natural next steps.

**Hardware Acceleration.** While QFP is hardware-agnostic, specialized accelerators for operator execution could improve performance without altering safety semantics.

**Formal Verification.** Further formalization of Safety Kernel algorithms and execution invariants could enable machine-checked proofs of compliance.

**Interoperable Policy Standards.** Standardized representations of policy artifacts may facilitate interoperability across organizations and jurisdictions.

**Expanded Compliance Proofs.** Richer zero-knowledge predicates could support more nuanced regulatory requirements without increased disclosure.

## 10.7 Relation to the Broader QVM Stack

Within the broader QVM ecosystem:

- QFM defines the mathematical language of computation,
- QFP governs whether and how that language may execute,
- QVM provides scalable deployment and orchestration.

Each layer addresses a distinct concern while remaining compatible with the others.

## 10.8 Concluding Perspective

The Quantum Field Processor is not proposed as a universal solution, but as a principled architectural response to a specific problem: executing powerful, global computations safely and accountably.

By embedding regulation into execution rather than applying it after the fact, QFP provides a path toward computational systems that are both capable and governable.

Future work will determine how far this approach can be extended, but the architecture presented here establishes a clear and technically grounded starting point.